

Debugging Malloc Lab

Introduction

The usual implementation of malloc and free are unforgiving to errors in their callers' code, including cases where the programmer overflows an array, forgets to free memory, or frees a memory block twice. This often does not affect the program immediately, waiting until the corrupted memory is used later (in the case of overwrites) or gradually accumulating allocated but unused blocks. Thus, debugging can be extremely difficult.

In this assignment, you will write a wrapper for the malloc package that will catch errors in the code that calls malloc and free. The skills you will have learned upon the completion of this exercise are pointer arithmetic and a greater understanding of the consequences of subtle memory mistakes.

Logistics

Unzip **debugging_malloc.zip** into an empty directory. The files contained are as follows:

File(s):	Function:
debugmalloc.c	Contains the implementation of the three functions you will be writing. This is the one file you will be editing and handing in.
debugmalloc.h	Contains the declaration of the functions, as well as the macros that will call them.
driver.c	Contains main procedure and the code that will be calling the functions in the malloc package
dmhelper.c, dmhelper.h	Contain the helper functions and macros that you will be calling from your code
grader.pl	Perl script that runs your code for the various test cases and gives you feedback based on your current code
debugmalloc.dsp	Exercise 3 project file
debugmalloc.dsw	Exercise 3 workspace file
tailor.h, getopt.c, getopt.h	Tools that are used only by the driver program for I/O purposes. You will not need to know what the code in these files do.
Others	Required by Visual C++. You do not need to understand their purpose

Specification

Programs that use this package will call the macro `MALLOC` and `FREE`. `MALLOC` and `FREE` are used exactly the same way as the `malloc()` and `free()` functions in the standard C `malloc` package. That is, the line

```
void *ptr = MALLOC ( n );
```

will allocate a **payload** of at least `n` bytes, and `ptr` will point to the front of this block. The line

```
FREE(ptr);
```

will cause the payload pointed to by `ptr` to be deallocated and become available for later use. The macros are defined as follows:

```
#define MALLOC(s) MyMalloc(s, __FILE__, __LINE__)  
#define FREE(p) MyFree(p, __FILE__, __LINE__)
```

The `__FILE__` macro resolves to the filename and `__LINE__` resolves to the current line number.

The `debugmalloc.c` file contains three functions that you are required to implement, as shown:

```
void *MyMalloc(size_t size, char *filename, int linenumber);  
void MyFree(void *ptr, char *filename, int linenumber);  
int AllocatedSize();
```

Using the macros above allow `MyMalloc` and `MyFree` to be called with the filename and line number of the actual `MALLOC` and `FREE` calls, while retaining the same form as the usual `malloc` package. By default, `MyMalloc()` and `MyFree()` simply call `malloc()` and `free()`, respectively, and return immediately. `AllocatedSize()` should return the number of bytes currently allocated by the user: the sum of the requested bytes through `MALLOC` minus the bytes freed using `FREE`. By default, it simply returns 0 and thus is unimplemented. The definitions are shown below:

```
void *MyMalloc(size_t size, char *filename, int linenumber) {  
  
    return (malloc(size));  
}
```

```

void MyFree(void *ptr, char *filename, int linenumber) {

    free(ptr);
}

int AllocatedSize() {

    return 0;
}

```

Your job is to modify these functions so that they will catch a number of errors that will be described in the next section.

There are also two optional functions in the debugmalloc.c file that you can implement:


```
void PrintAllocatedBlocks(); int HeapCheck();
```

PrintAllocatedBlocks should print out information about all currently allocated blocks.

HeapCheck should check all the blocks for possible memory overwrites.

Implementation Details

To catch the errors, you will allocate a slightly larger amount of space and insert a header and a footer around the "requested payload". MyMalloc() will insert information into this area, and MyFree() will check to see if the information has not changed. The organization of the complete memory block is as shown below:



Note: MyMalloc() returns a pointer to the payload, *not* the beginning of the whole block. Also, the ptr parameter passed into MyFree(void *ptr) will point to the payload, *not* the beginning of the block.

Information that you might want to store in this extra (header, footer) area include:

- a "fence" immediately around the requested payload with a known value like 0xCCDEADCC, so that you can check if it has been changed when the block is freed.
- the size of the block
- a checksum for the header to ensure that it has not been corrupted (A checksum of a sequence of bits is calculated by counting the number of "1" bits in the stream. For example, the checksum for "1000100010001000" is 4. It is a simple error detection mechanism.)
- the filename and line number of the MALLOC() call

The errors that can occur are:

- Error #1: Writing past the beginning of the user's block (through the fence)
- Error #2: Writing past the end of the user's block (through the fence)
- Error #3: Corrupting the header information
- Error #4: Attempting to free an unallocated or already-freed block
- Error #5: Memory leak detection (user can use ALLOCATEDSIZE to check for leaks at the end of the program)

To report the first four errors, call one of these two functions:

```
void error(int errorcode, char *filename, int linenumber);
    errorcode is the number assigned to the error as
    stated above. filename and linenumber contain the
    filename and line number of the line (thefree call) in
    which the error is invoked. For example, call error(2,
    filename, linenumber) if you come across a
    situation where the footer fence has been changed.
void errorfl(int errorcode, char *filename_malloc, int
linenumber_malloc, char *filename_free, int linenumber_free);
    This is the same as the error(), except there are
    two sets of filenames and line numbers, one for the
    statement in which the block was malloc'd, and the
    other for the statement in which the block was
    free'd (and the error was invoked).
```

The fact that MyMalloc() and MyFree() are given the filename and line number of the MALLOC() and FREE() call can prove to be very useful when you are reporting errors. The more information you print out, the easier it will be for the programmer to locate the error. Use errorfl() instead of error() whenever possible. errorfl() obviously cannot be used on situations where FREE() is called on an unallocated block, since it was not ever MALLOC'd.

Note: You will only be reporting errors from MyFree(). None of the errors can be caught in MyMalloc()

In the case of memory leaks, the driver program will call `AllocatedSize()`, and the grader will look at its return value and possible output. `AllocatedSize()` should return the number of bytes currently allocated from `MALLOC` and `FREE` calls. For example, the code segment:

```
void *ptr1 = MALLOC(10), *ptr2 = MALLOC(8); FREE(ptr2);  
printf("%d\n", AllocatedSize());
```

should print out "10".

Once you have gotten to the point where you can catch all of the errors, you can go an optional step further and create a global list of allocated blocks. This will allow you to perform analysis of memory leaks and currently allocated memory. You can implement the void `PrintAllocatedBlocks()` function, which prints out the filename and line number where all currently allocated blocks were `MALLOC()`'d. A macro is provided for you to use to print out information about a single block in a readable and gradeable format:

```
PRINTBLOCK(int size, char *filename, int linenum)
```

Also, you can implement the int `HeapCheck()` function. This should check all of the currently allocated blocks and return -1 if there is an error and 0 if all blocks are valid. In addition, it should print out the information about all of the corrupted blocks, using the macro `#define PRINTERRORError(int errorcode, char *filename, int linenum)`, with `errorcode` equal to the error number (according to the list described earlier) the block has gone through.

You may find that this global list can also allow you to be more specific in your error messages, as it is otherwise difficult to determine the difference between an overwrite of a non-payload area and an attempted `FREE()` of an unallocated block.

Evaluation

You are given 7 test cases to work with, plus 1 extra for testing a global list. You can type `"debugmalloc -t n"` to run the n-th test. You can see the code that is being run in `driver.c`. There are a total of 100 possible points.

Here is a rundown of the test cases and desired output (do not worry about the path of the filename):

Test case #1	
Code	<code>char *str = (char *) MALLOC(12); strcpy(str, "123456789");</code>

	FREE(str); printf("Size: %d\n", AllocatedSize()); PrintAllocatedBlocks();
Error #	None
Correct Output	Size: 0
Points worth	10
Details	10 points for not reporting an error and returning 0 in AllocatedSize()

Test case #2	
Code	char *str = (char *) MALLOC(8); strcpy(str, "12345678"); FREE(str);
Error #	2
Correct Output	Error: Ending edge of the payload has been overwritten. in block allocated at driver.c, line 21 and freed at driver.c, line 23
Points worth	15
Details	6 pts for catching error 3 pts for printing the filename/line numbers 6 pts for correct error message

Test case #3	
Code	char *str = (char *) MALLOC(2); strcpy(str, "12"); FREE(str);
Error #	2
Correct Output	Error: Ending edge of the payload has been overwritten. in block allocated at driver.c, line 28 and freed at driver.c, line 30
Points worth	15
Details	6 pts for catching error 3 pts for printing the filename/line numbers 6 pts for correct error message

Test case #4

Code	<code>void *ptr = MALLOC(4); *ptr2 = MALLOC(6); FREE(ptr); printf("Size: %d\n", AllocatedSize()); PrintAllocatedBlocks();</code>
Error #	None
Correct Output	Size: 6 Currently allocated blocks: 6 bytes, created at driver.c, line 34
Points worth	15
Details	15 pts for not reporting an error and returning 6 from AllocatedSize Extra for printing out the extra block

Test case #5	
Code	<code>void *ptr = MALLOC(4); FREE(ptr); FREE(ptr);</code>
Error #	4
Correct Output	Error: Attempting to free an unallocated block. in block freed at driver.c, line 43
Points worth	15
Details	15 pts for catching error Extra for correct error message

Test case #6	
Code	<code>char *ptr = (char *) MALLOC(4); *((int *) (ptr - 8)) = 8 + (1 << 31); FREE(ptr);</code>
Error #	1 or 3
Correct Output	Error: Header has been corrupted. or Error: Starting edge of the payload has been overwritten. in block allocated at driver.c, line 47 and freed at driver.c, line 49
Points worth	15
Details	9 pts for catching error 6 pts for a correct error message

Test case #7	
Code	<code>char ptr[5]; FREE(ptr);</code>

Error #	4
Correct Output	Error: Attempting to free an unallocated block. in block freed at driver.c, line 54
Points worth	15
Details	15 pts for recognizing error Extra for printing correct error message

Test case #8 (Optional)	
Code	int i; int *intptr = (int *) MALLOC(6); char *str = (char *) MALLOC(12); for(i = 0; i < 6; i++) { intptr[i] = i; } if (HeapCheck() == -1) { printf("\nCaught Errors\n"); }
Error #	None
Correct Output	Error: Ending edge of the payload has been overwritten. Invalid block created at driver.c, line 59 Caught Errors
Points worth	Extra
Details	"Caught Errors" indicates that the HeapCheck() function worked correctly. Extra points possible.

Your instructor may give you extra credit for implementing a global list and the PrintAllocatedBlocks() and HeapCheck() functions.